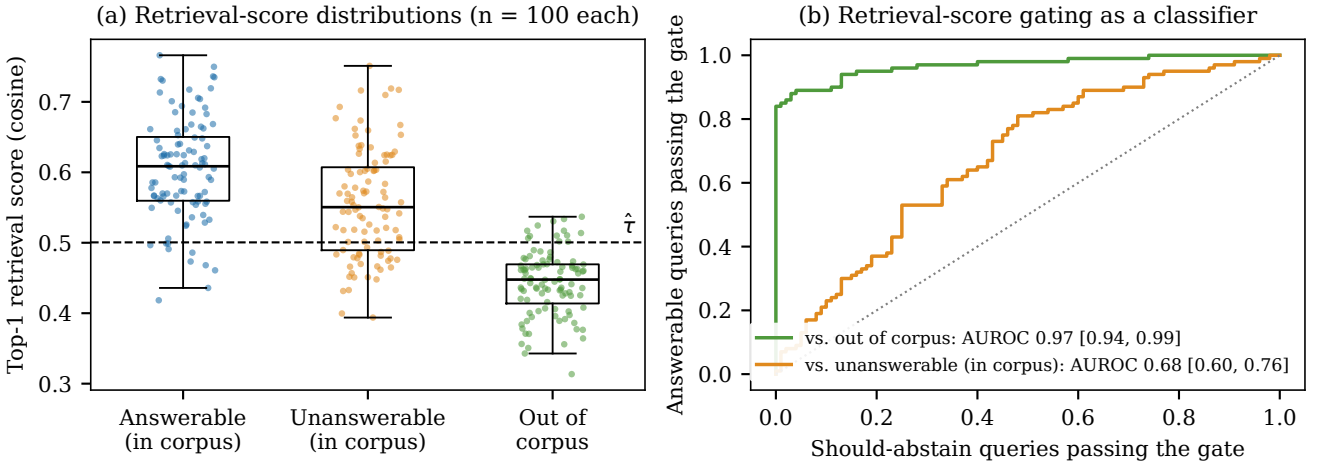
(a) Retrieval-score distributions (n = 100 each)
Top-1 retrieval score (cosine)
0.7
0.6
0.5
0.4
0.3
$\hat{\tau}$
Answerable (in corpus)
Unanswerable (in corpus)
Out of corpus
(b) Retrieval-score gating as a classifier
Answerable queries passing the gate
1.0
0.8
0.6
0.4
0.2
0.0
vs. out of corpus: AUROC 0.97 [0.94, 0.99]
vs. unanswerable (in corpus): AUROC 0.68 [0.60, 0.76]
0.0
0.2
0.4
0.6
0.8
1.0
Should-abstain queries passing the gate